

CMPE 220

Class 10 – Operating Systems - part 1

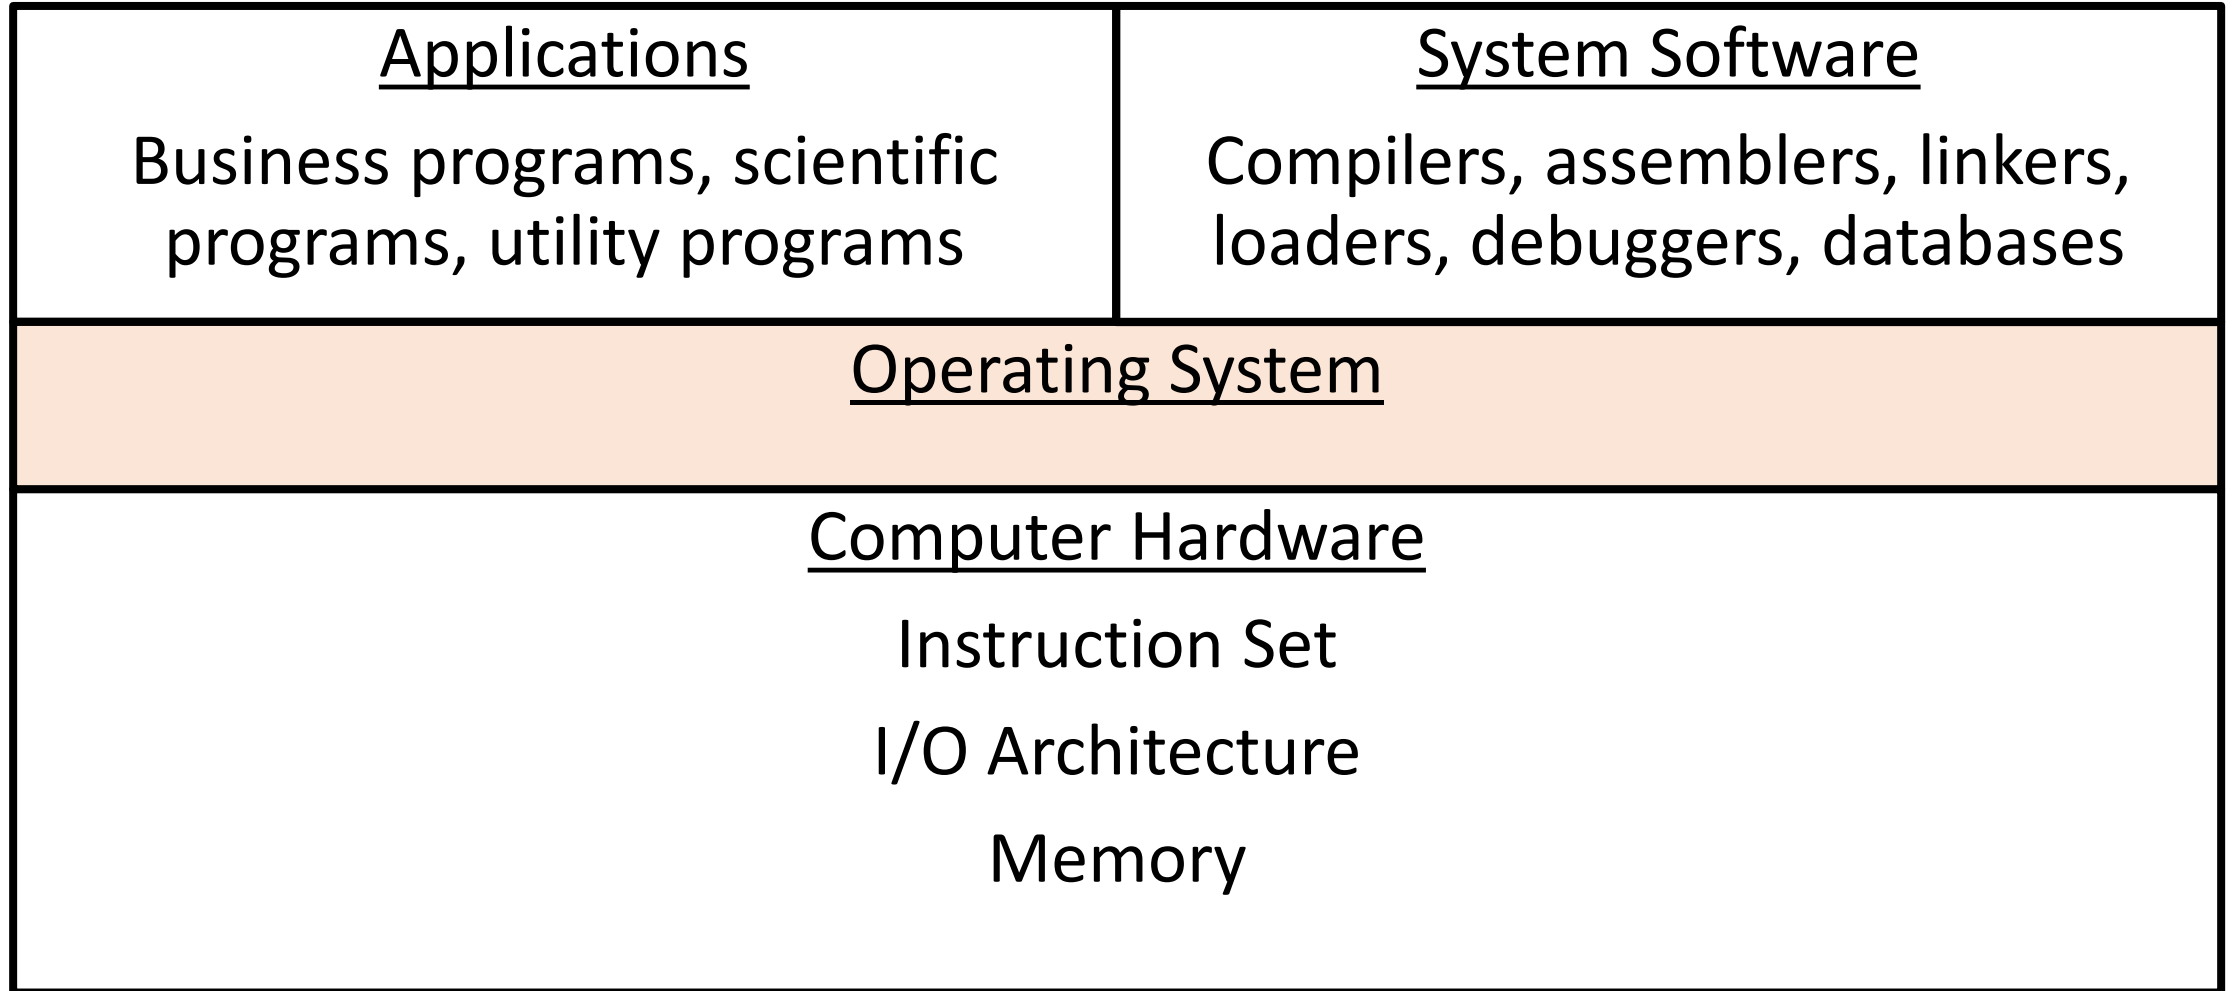
What is an Operating System?

Wikipedia: An operating system (OS) is system software that manages computer hardware, software resources, and provides common services for computer programs.

TechTerms: An operating system, or "OS," is software that communicates with the hardware and allows other programs to run. It is comprised of system software, or the fundamental files your computer needs to boot up and function.

HowToGeek: An operating system is the primary software that manages all the hardware and other software on a computer. The operating system, also known as an "OS," interfaces with the computer's hardware and provides services that applications can use.

What is an Operating System?

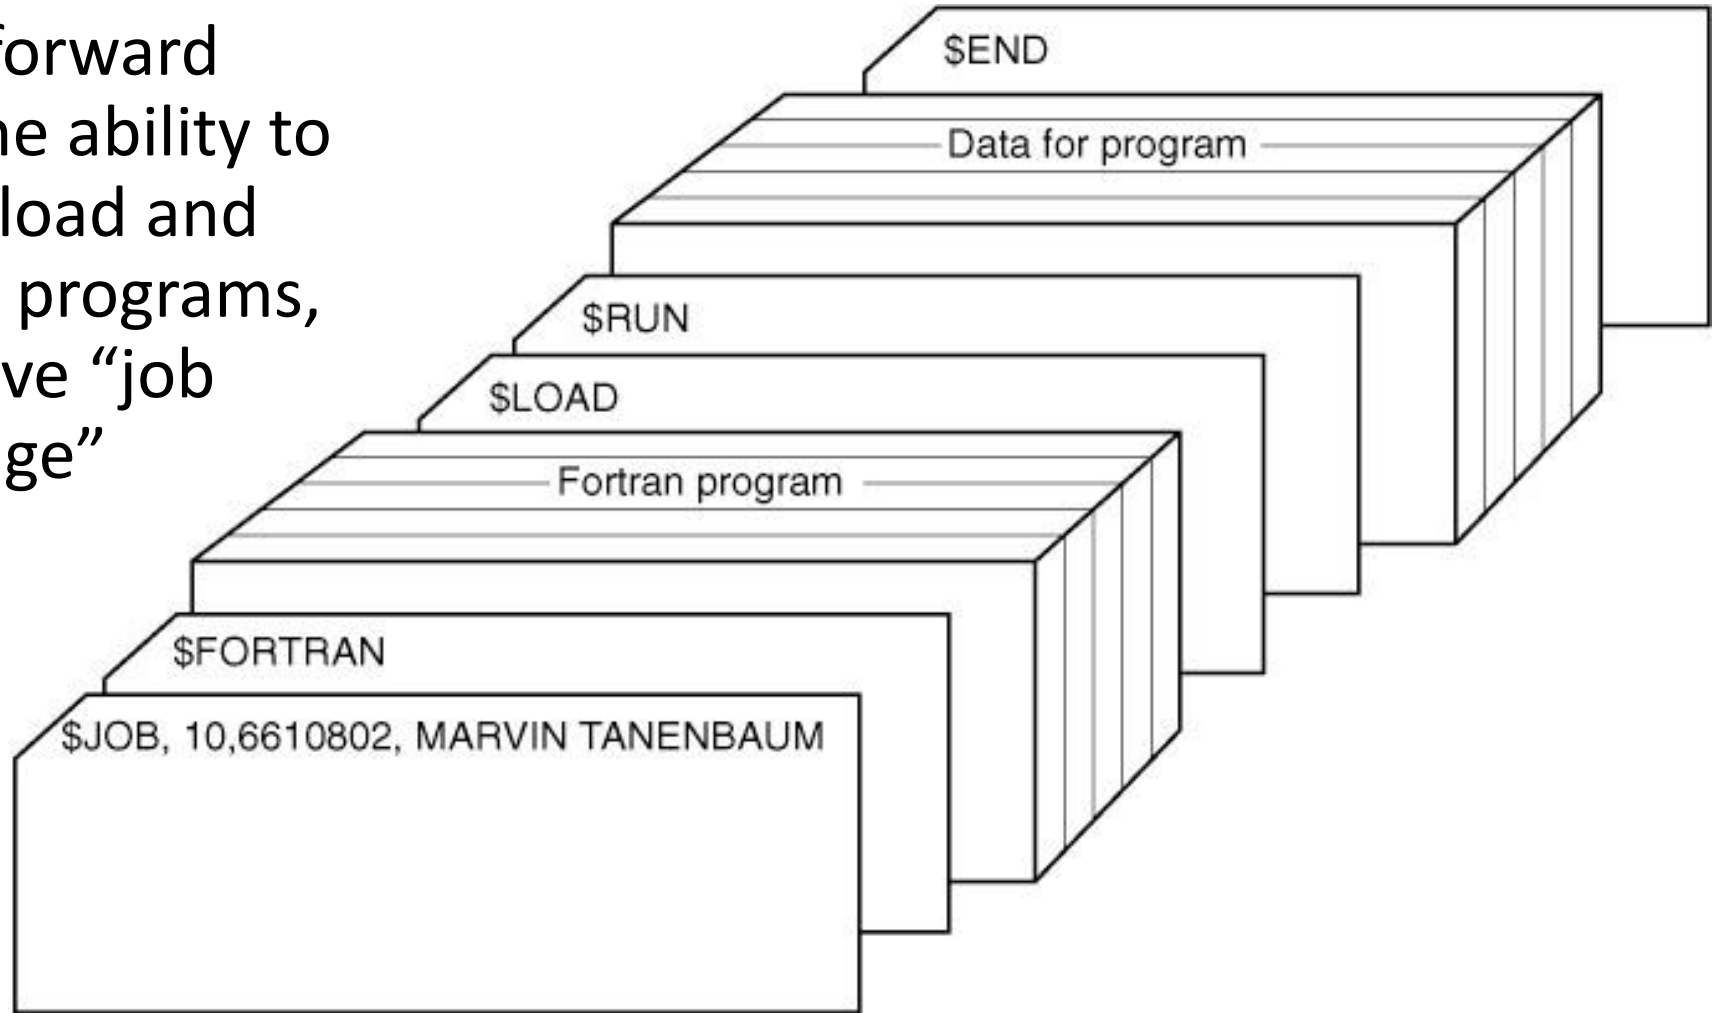


History

- The first computers (1940s) did not have an operating system.
- Computers had a primitive *absolute loader* – often stored in some form of non-volatile memory
- Programs were loaded – typically from punched cards – and had complete control over the hardware
 - Programmers were responsible for memory management and I/O
- When the program was done (or crashed), the next user would take over the computer and load their program
- Microcomputers in the 1970s followed the same path, except programs were usually loaded from punched paper tape

History: Job Control Languages (JCL)

- The first step forward (1950s) was the ability to automatically load and run a series of programs, using a primitive “job control language” (JCL)



History: Job Control Programs

- This required a small “job control program” to remain **resident** in memory.
- This was a tradeoff. It used (precious) memory, but it made more efficient use of the computer
 - *Batch* processing
 - Moved on quickly when programs didn't work
 - Did not require the programmer to be present
- This began a slippery slope
 - Saving work from programmers by adding functions to the resident job control program. Programming became simpler, but the control program got bigger.


A Typical Job

- A Fortran Compiler
 - 3 ½ boxes of punched cards
 - Each box = 2000 cards (about ten pounds each)
 - A 35 pound program!



Late 1950s: Recognizable Operating Systems

- IBM 1401 (1959)




The IBM 1401: the Model T of Computing


- 1959-1971
- Inexpensive
- Decimal (BCD) Arithmetic
- High Sales Volume: over 12,000 sold
 - By the mid-1960s, almost half the computers in the world were 1401s
- There is a fully restored and working IBM 1401 at the Computer History Museum in Mountain View

Interested in Early IBM History?

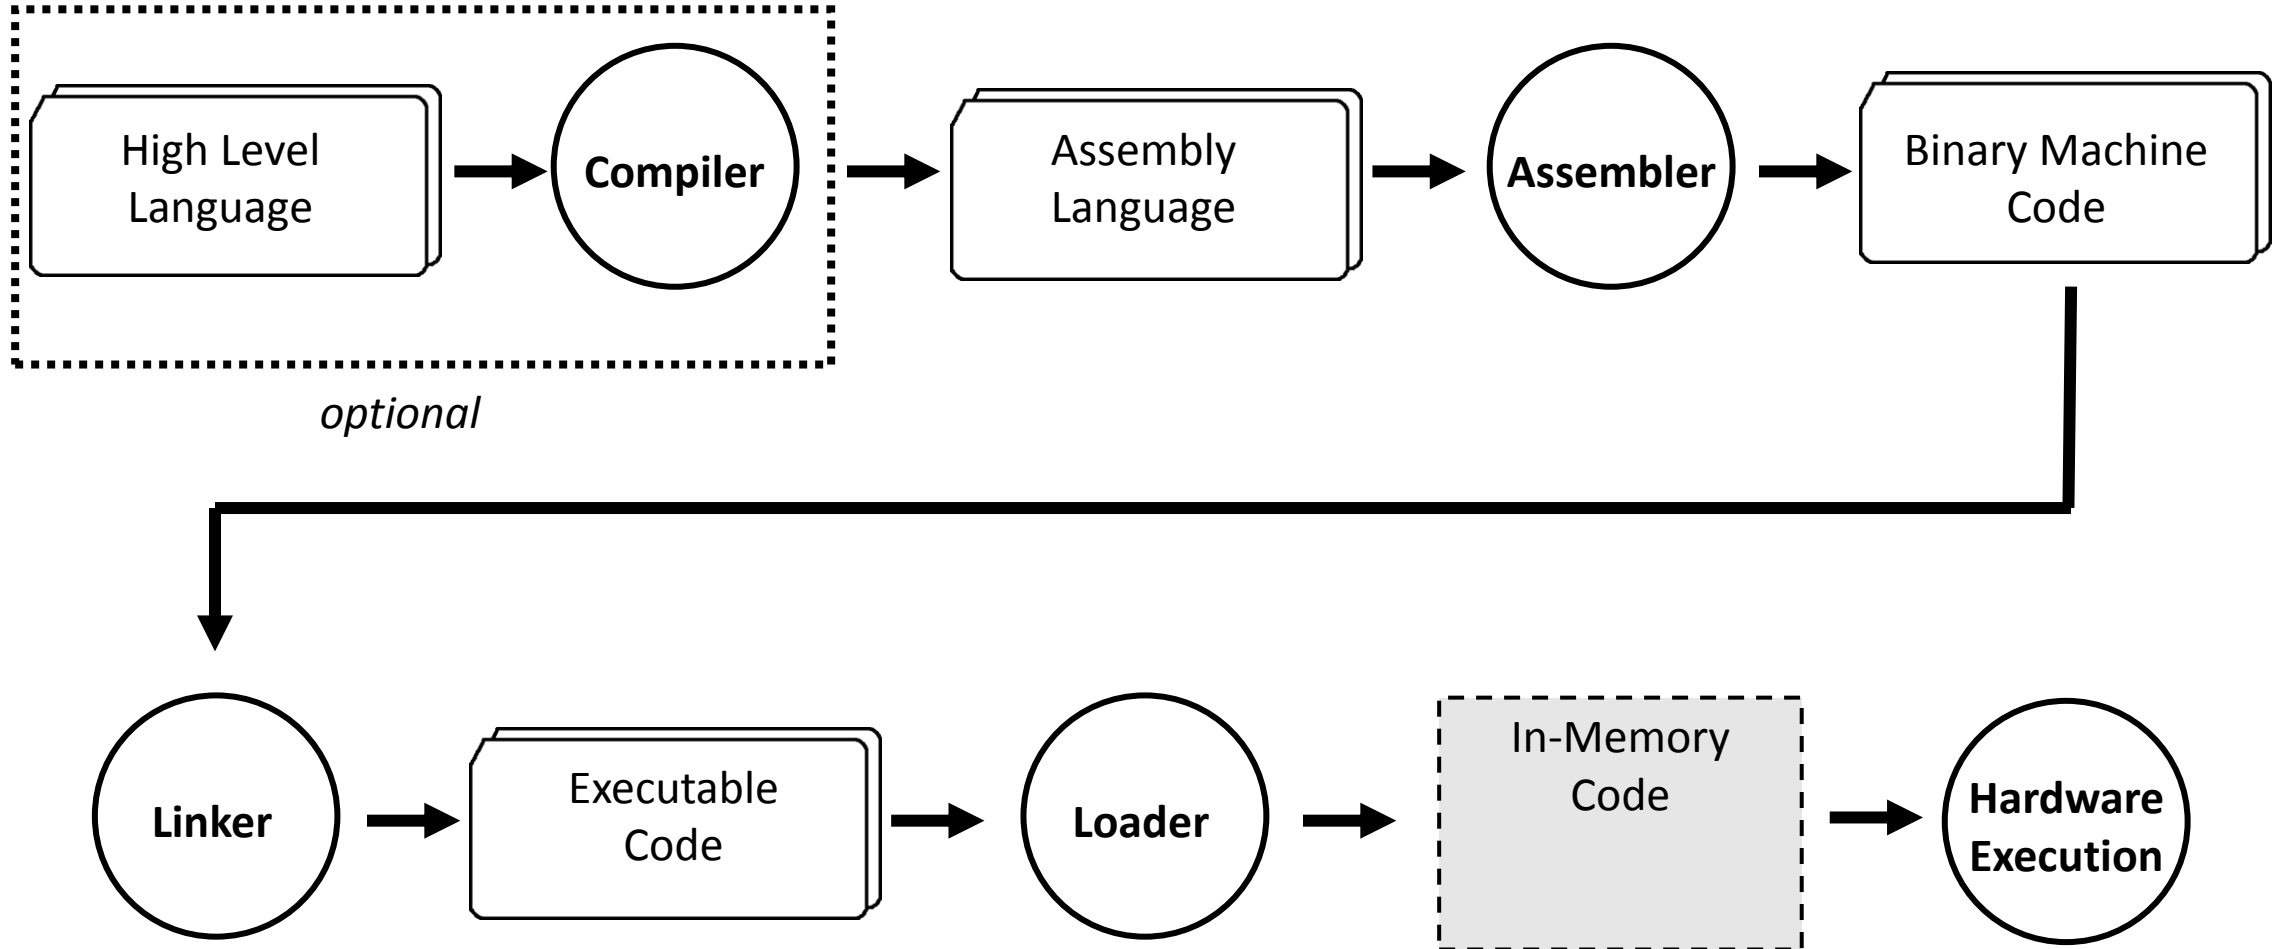
- Light Blue: An Entry Level and Mid Management Perspective of IBM's Evolution Through the Golden Years
 - Justin “Jud” McCarty
 - June, 2020



Building Software



Building Software: 1950s



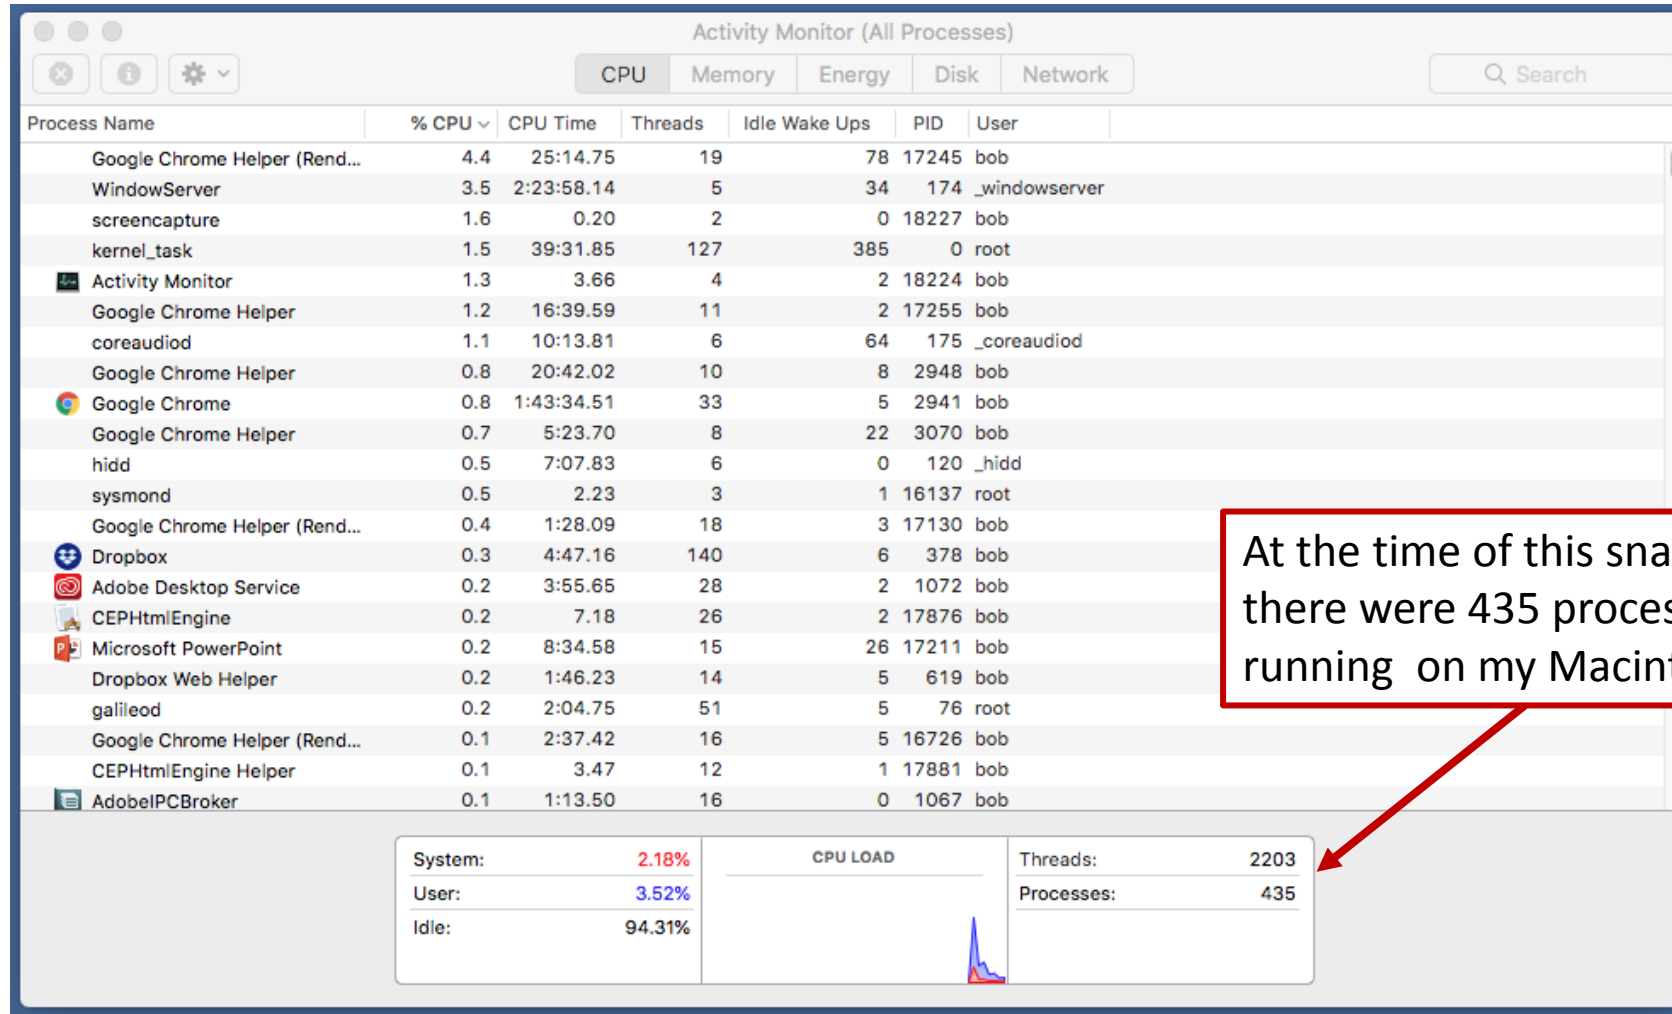
What Does a Modern Operating System Do?

1. Process Management
 - Interprocess Communications
2. Input / Output (I/O) Management
3. Memory Management
4. File System Management
5. System Functions and Kernel Mode
6. User Interaction – (maybe)

(1) Process Management

- A modern computer runs many programs at the same time
- Each instance of a running program is called a *process*
 - A program is just code
 - A process is code, data, and state information *running on a computer*
- Each process has its own address space... in effect, it behaves as if it is the only program running
- *It's important to understand that the operating system itself runs as one or more of these processes*

Processes – Mac OS: Activity Monitor



Processes – POSIX: ps -axl command

```
bob$ ps -axl
```

UID	PID	PPID	F	CPU	PRI	NI	SZ	RSS	WCHAN	S	ADDR	TTY	TIME	CMD
0	1	0	4004	0	37	0	4330580	14228	-	Ss	0 ??		5:13.39	/sbin/launch
0	59	1	4004	0	4	0	4306056	916	-	Ss	0 ??		0:11.62	/usr/sbin/sy
0	60	1	4004	0	37	0	4333164	10864	-	Ss	0 ??		0:14.06	/usr/libexec
0	64	1	4004	0	20	0	4296532	3304	-	Ss	0 ??		0:02.98	/System/Libr
0	65	1	4004	0	37	0	4339484	1904	-	Ss	0 ??		0:03.19	/usr/libexec
0	66	1	1004004	0	50	0	4356220	4368	-	Ss	0 ??		1:06.17	/System/Libr
0	70	1	4004	0	4	0	4311292	10024	-	Ss	0 ??		0:00.59	/Library/Fra
0	71	1	4004	0	4	0	4358892	6264	-	Ss	0 ??		0:10.08	/usr/sbin/sy
55	73	1	4004	0	4	0	4333692	4924	-	Ss	0 ??		0:02.59	/System/Libr
0	74	1	400c	0	37	0	4336288	5476	-	Ss	0 ??		0:11.68	/usr/libexec
0	75	1	4004	0	37	0	4331460	4864	-	Ss	0 ??		0:38.00	/System/Libr
0	76	1	4004	0	20	0	8929896	15168	-	Ss	0 ??		2:04.89	/Library/App
0	79	1	4004	0	37	0	4413840	18912	-	Ss	0 ??		0:45.90	/usr/libexec
0	83	1	4004	0	37	10	4336048	6708	-	SNs	0 ??		0:02.35	/usr/libexec
0	84	1	1004004	0	50	0	4504488	37392	-	Ss	0 ??		48:37.12	/System/Libr
0	87	1	4004	0	4	0	4304388	164	-	Ss	0 ??		0:00.02	firmwaresync
0	91	1	4004	0	37	0	4330644	2864	-	Ss	0 ??		0:03.30	/usr/libexec
0	93	1	4004	0	4	0	4348432	23776	-	Ss	0 ??		0:34.61	/usr/libexec
0	94	1	4004	0	4	0	4305144	14416	-	Ss	0 ??		0:00.42	/System/Libr
0	98	1	4004	0	37	0	4345252	8284	-	Ss	0 ??		0:43.12	/usr/libexec
0	101	1	4004	0	4	0	4331928	2632	-	Ss	0 ??		0:00.25	/System/Libr
0	102	1	4004	0	4	0	4338264	10756	-	Ss	0 ??		0:22.02	/System/Libr
0	104	1	4004	0	37	0	4338644	5944	-	Ss	0 ??		0:38.26	/usr/sbin/se
205	105	1	4004	0	4	0	4337968	6628	-	Ss	0 ??		0:08.41	/usr/libexec
244	108	1	4104	0	20	0	4305248	3104	-	Ss	0 ??		0:00.07	/usr/libexec
0	109	1	4004	0	20	0	4304884	188	-	Ss	0 ??		0:00.02	autofs
0	111	1	4004	0	4	0	4336212	7104	-	Ss	0 ??		0:05.38	/usr/libexec
74	113	1	4044	0	31	0	5372540	1544	-	Ss	0 ??		0:21.59	/usr/local/m
0	114	1	1004004	0	4	0	4370036	3332	-	Ss	0 ??		0:00.38	/System/Libr
504	115	1	80004104	0	63	0	4466888	39956	-	Ss	0 ??		0:25.84	/System/Libr
0	116	1	4004	0	37	0	4304496	1936	-	Ss	0 ??		0:00.13	/System/Libr
0	117	1	4004	0	37	0	4296464	240	-	Ss	0 ??		0:00.01	/usr/sbin/Ke

How Processes Are Created

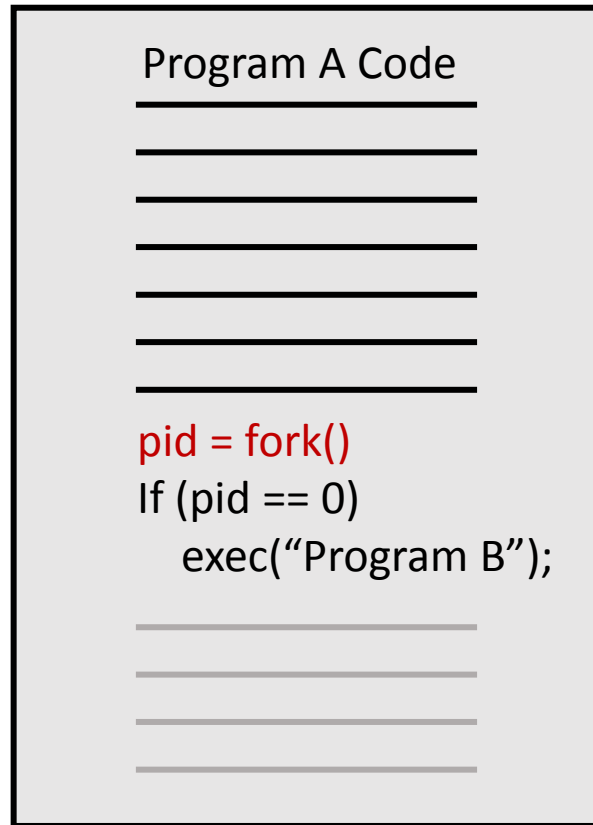
- A system function call is required to create a new process
- A system function call is required to invoke the loader, load a program, and start execution
- There may be two separate calls, or a single call to do both
- In POSIX systems:
 - The `fork()` system call duplicates the currently running process – leaving TWO processes running the identical code, with identical states
 - The `exec()` system call loads an executable into the current process
 - For one program to “launch” another, it first calls `fork()`, and then the child process calls `exec()`

Fork() and Exec()


```
If ((pid = fork()) < 0) { // After this call there should be 2 processes running
    printf("Fork failed\n");
    exit(-1);
}
If (pid == 0) { // This is the child process
    exec( "/users/robert/fibonnacci" ); // Invoke the loader
    printf("Exec failed\n");
    exit(-1);
}
// Parent process continues
printf("Process ID of child is %d\n", pid);
```

Fork() and Exec()

Process One



Process Two



How Processes are Terminated

- The program may call a system function, such as the POSIX *exit()* function, that does some cleanup and deletes the process
- When a program terminates without calling *exit()*, the system automatically executes some code that does the same thing
 - Alluded to in earlier lectures
- With appropriate permissions, one process may force the deletion of another process by calling a system function such as the POSIX *kill()* function
- The system may also have security features that will automatically kill processes that use too much CPU, memory, etc.


Interrupts

- A multiprocessing operating system relies on *interrupts*
- An interrupt is a signal sent to the processor that interrupts the current process
- It may be generated by a hardware device or a software program
- Types of interrupts
 - Timer
 - I/O state change (such as operation complete)
 - A signal from one process to another

What an Interrupt Does

Currently Executing Program

Interrupt Handler



Saving the Process State

- In order to allow us to resume execution of the current process, the interrupt handler must save the state (such as the registers) of the currently executing program in a data structure called a **Process Control Block (PCB)**
- There is a PCB for each process

Contents of a Process Control Block (PCB)

Process management	Memory management	File management
Registers Program counter Program status word Stack pointer Process state Time when process started CPU time used Children's CPU time Time of next alarm Message queue pointers Pending signal bits Process id Various flag bits	Pointer to text segment Pointer to data segment Pointer to bss segment Exit status Signal status Process id Parent process Process group Real uid Effective uid Real gid Effective gid Bit maps for signals Various flag bits	UMASK mask Root directory Working directory File descriptors Effective uid Effective gid System call parameters Various flag bits

Process Scheduling

- The operating system makes sure that every program gets some time to run.
- **Scheduler:** a component of the operating system that determines which process runs next
- **Dispatcher:** The task of switching control to another process is called *dispatching*, and the code that accomplishes this is called the *dispatcher*
- **Time Slicing:** The operating system uses a *timer interrupt* to periodically regain control so that it can switch from one process to another

Scheduling Algorithms

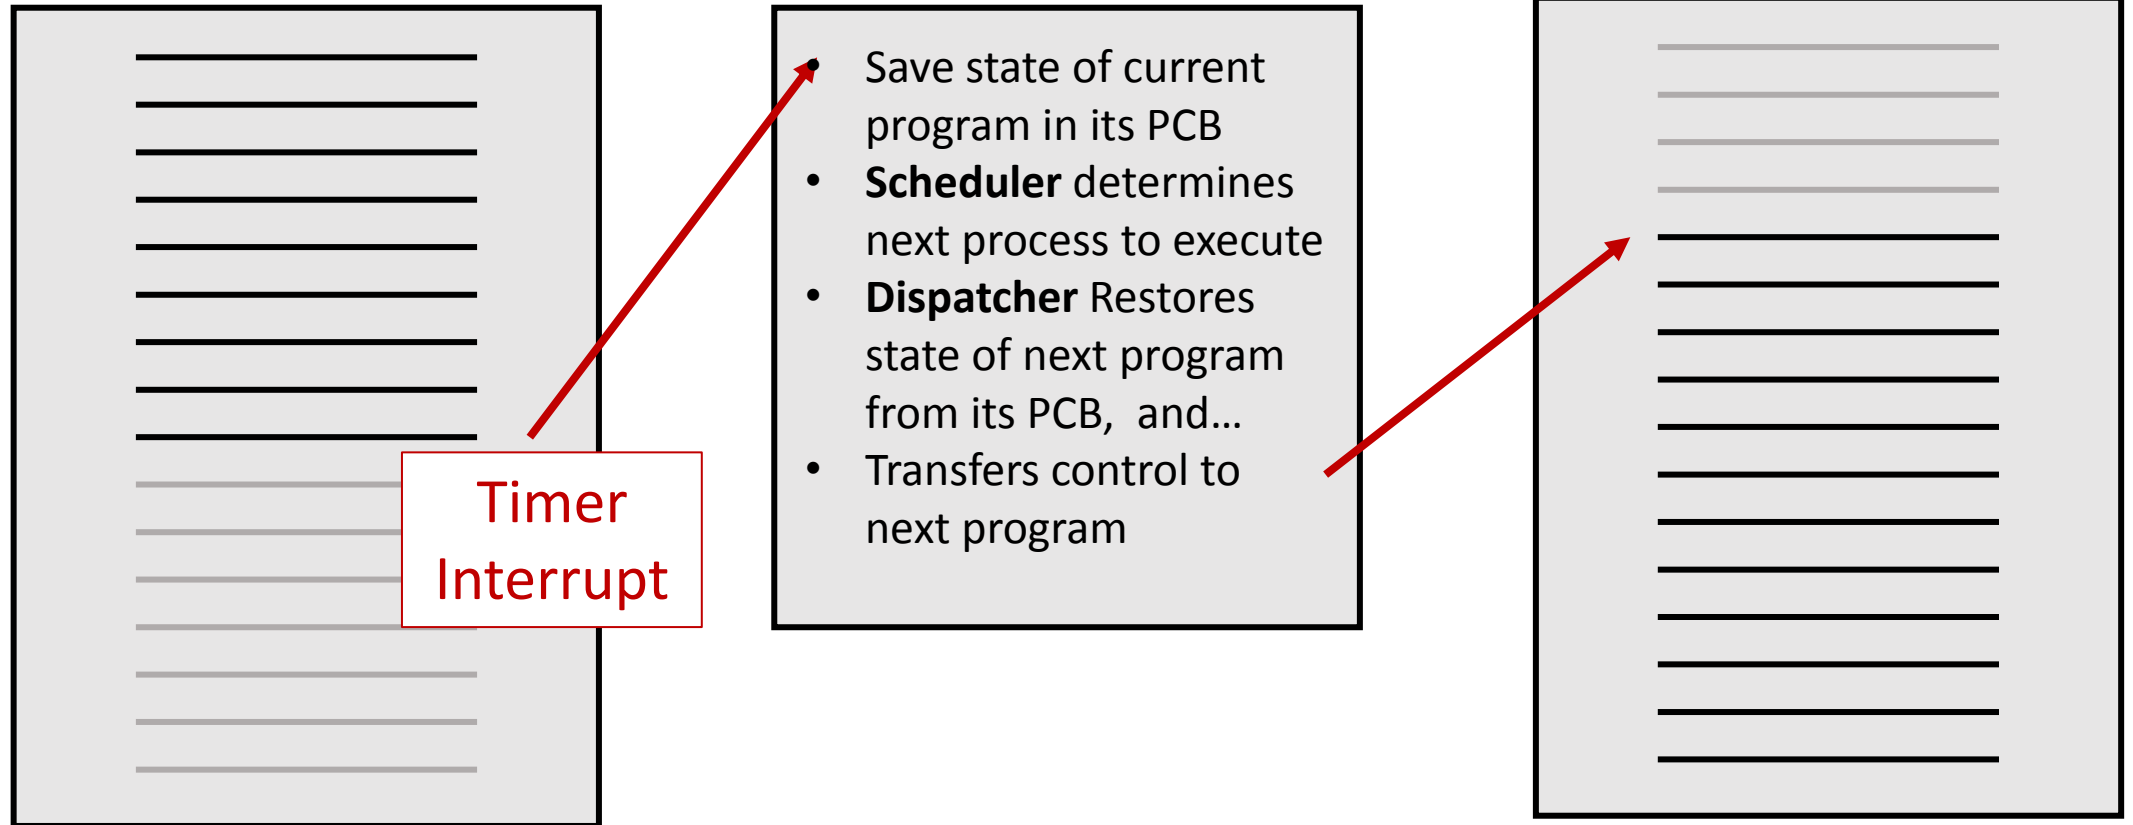
- **Round Robin:** All processes get an equal time-slice, in order
- **Priority Scheduling:** Some processes may be scheduled as higher priority, and get more or longer time slices
 - The original Lunar Lander had a single computer
 - The process that adjusted the attitude control jets ran at a higher priority than the process that updated the display
- **Adaptive Scheduling:** adjust scheduling based on process performance
 - A system may give less time to “CPU hogs”
- Many other algorithms are possible

Switching Processes

Currently Executing Process

Interrupt Handler

Next Process



Process Switching and I/O

- Typically, when a program starts an I/O operation, it can't proceed until the operation is complete
- In terms of processor speed, I/O operations take a really, really long time
- The system **scheduler** will not give control to a process that is waiting for I/O.
- We say that the process is *blocked* – or in an *I/O wait state*

What Happens When a Program Initiates I/O

- The program calls an operating system routine to start the I/O operation
- The operating system routine will:
 - Save the process state in its PCB
 - Set a flag in the process's PCB to indicate that the program is in an *I/O wait state*
 - Initiate the I/O operation
 - Call the **scheduler** to determine which process to execute next
 - Call the **dispatcher** to give control to that process

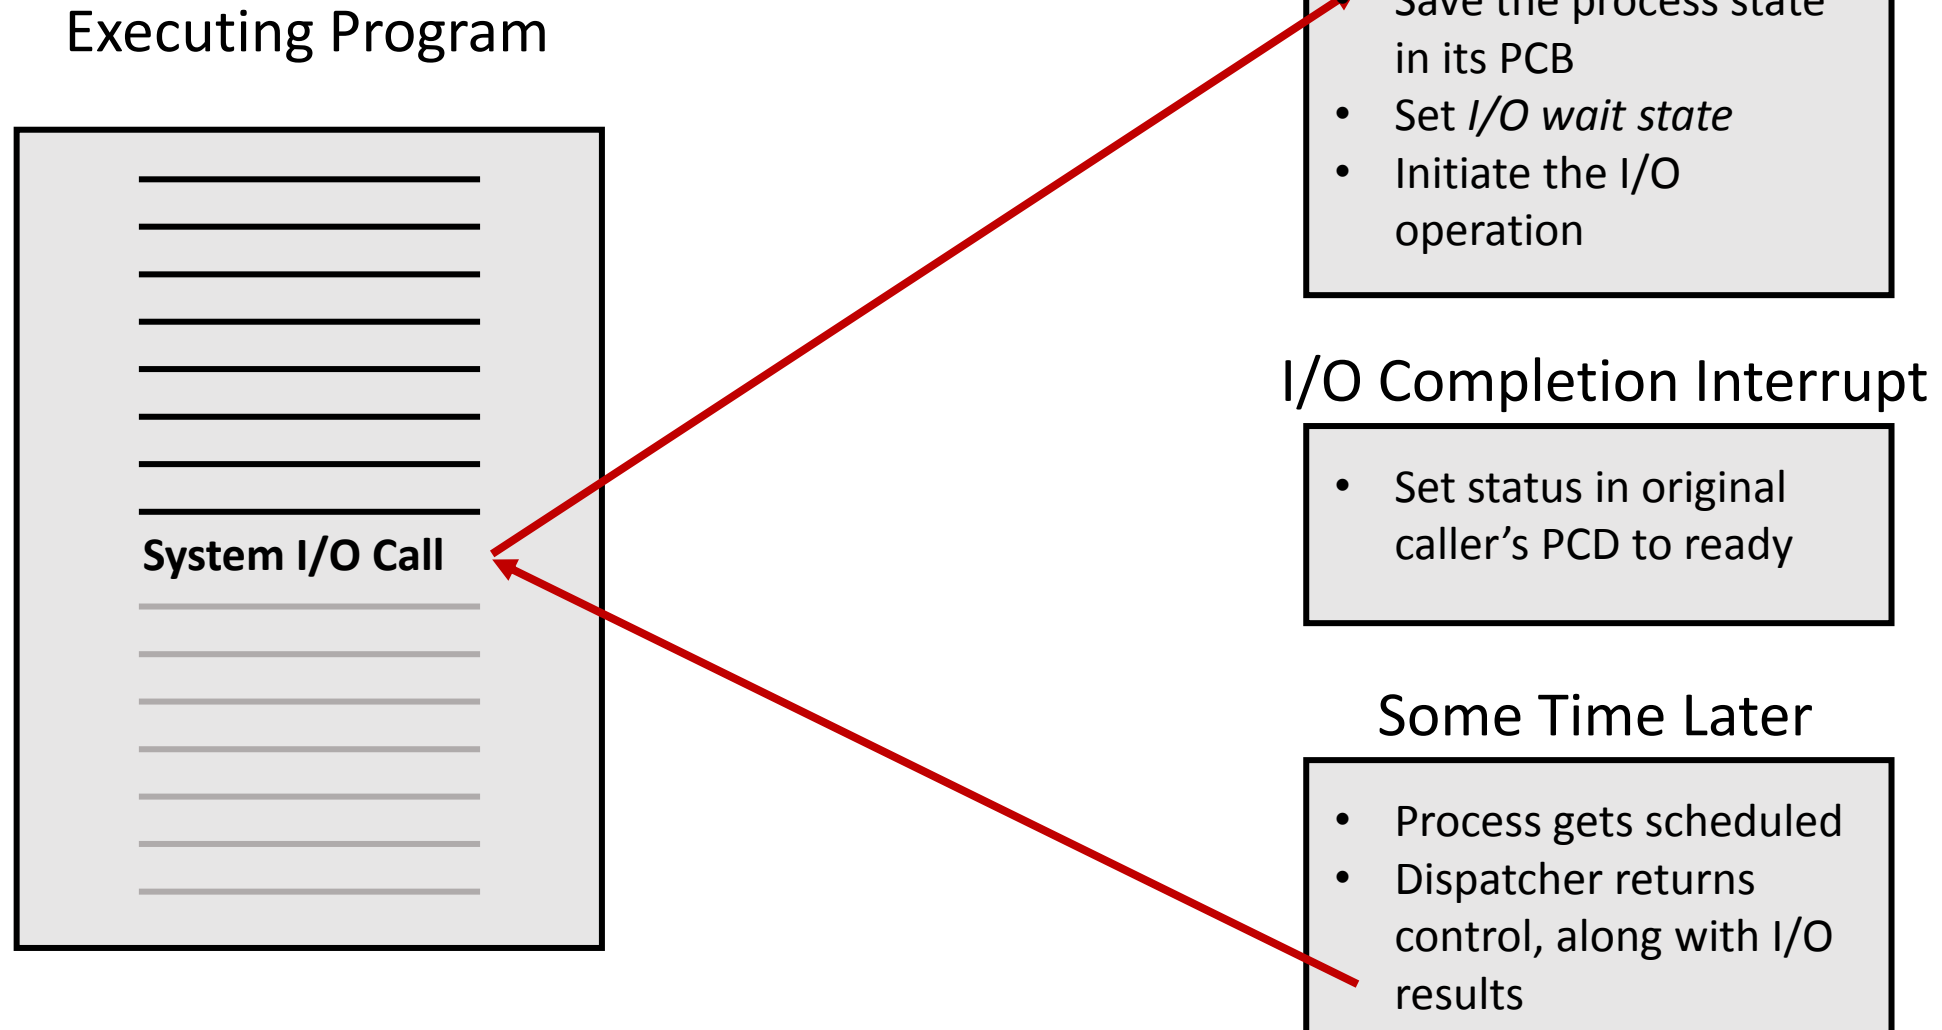
When an I/O Operation Completes

- A completed I/O operation generates an I/O Interrupt
- The interrupt handler gains control, and:
 - Saves the state of the currently executing process, whatever it may be
 - Determines which process initiated the I/O operation that just completed
 - Sets the status in that process's PCB to indicate it is *ready*
 - Calls the **scheduler** to determine which process to execute next
 - Calls the **dispatcher** to give control to that process
- Note that the “next process” *may or may not* be the process that initiated the I/O operation
- I/O completion simply makes that process *ready*, or eligible for scheduling

Returning Control to the I/O Caller

- Eventually, the dispatcher will return control to the process that initiated the I/O operation
- At that time, the system I/O function will return, passing back the results of the I/O operation

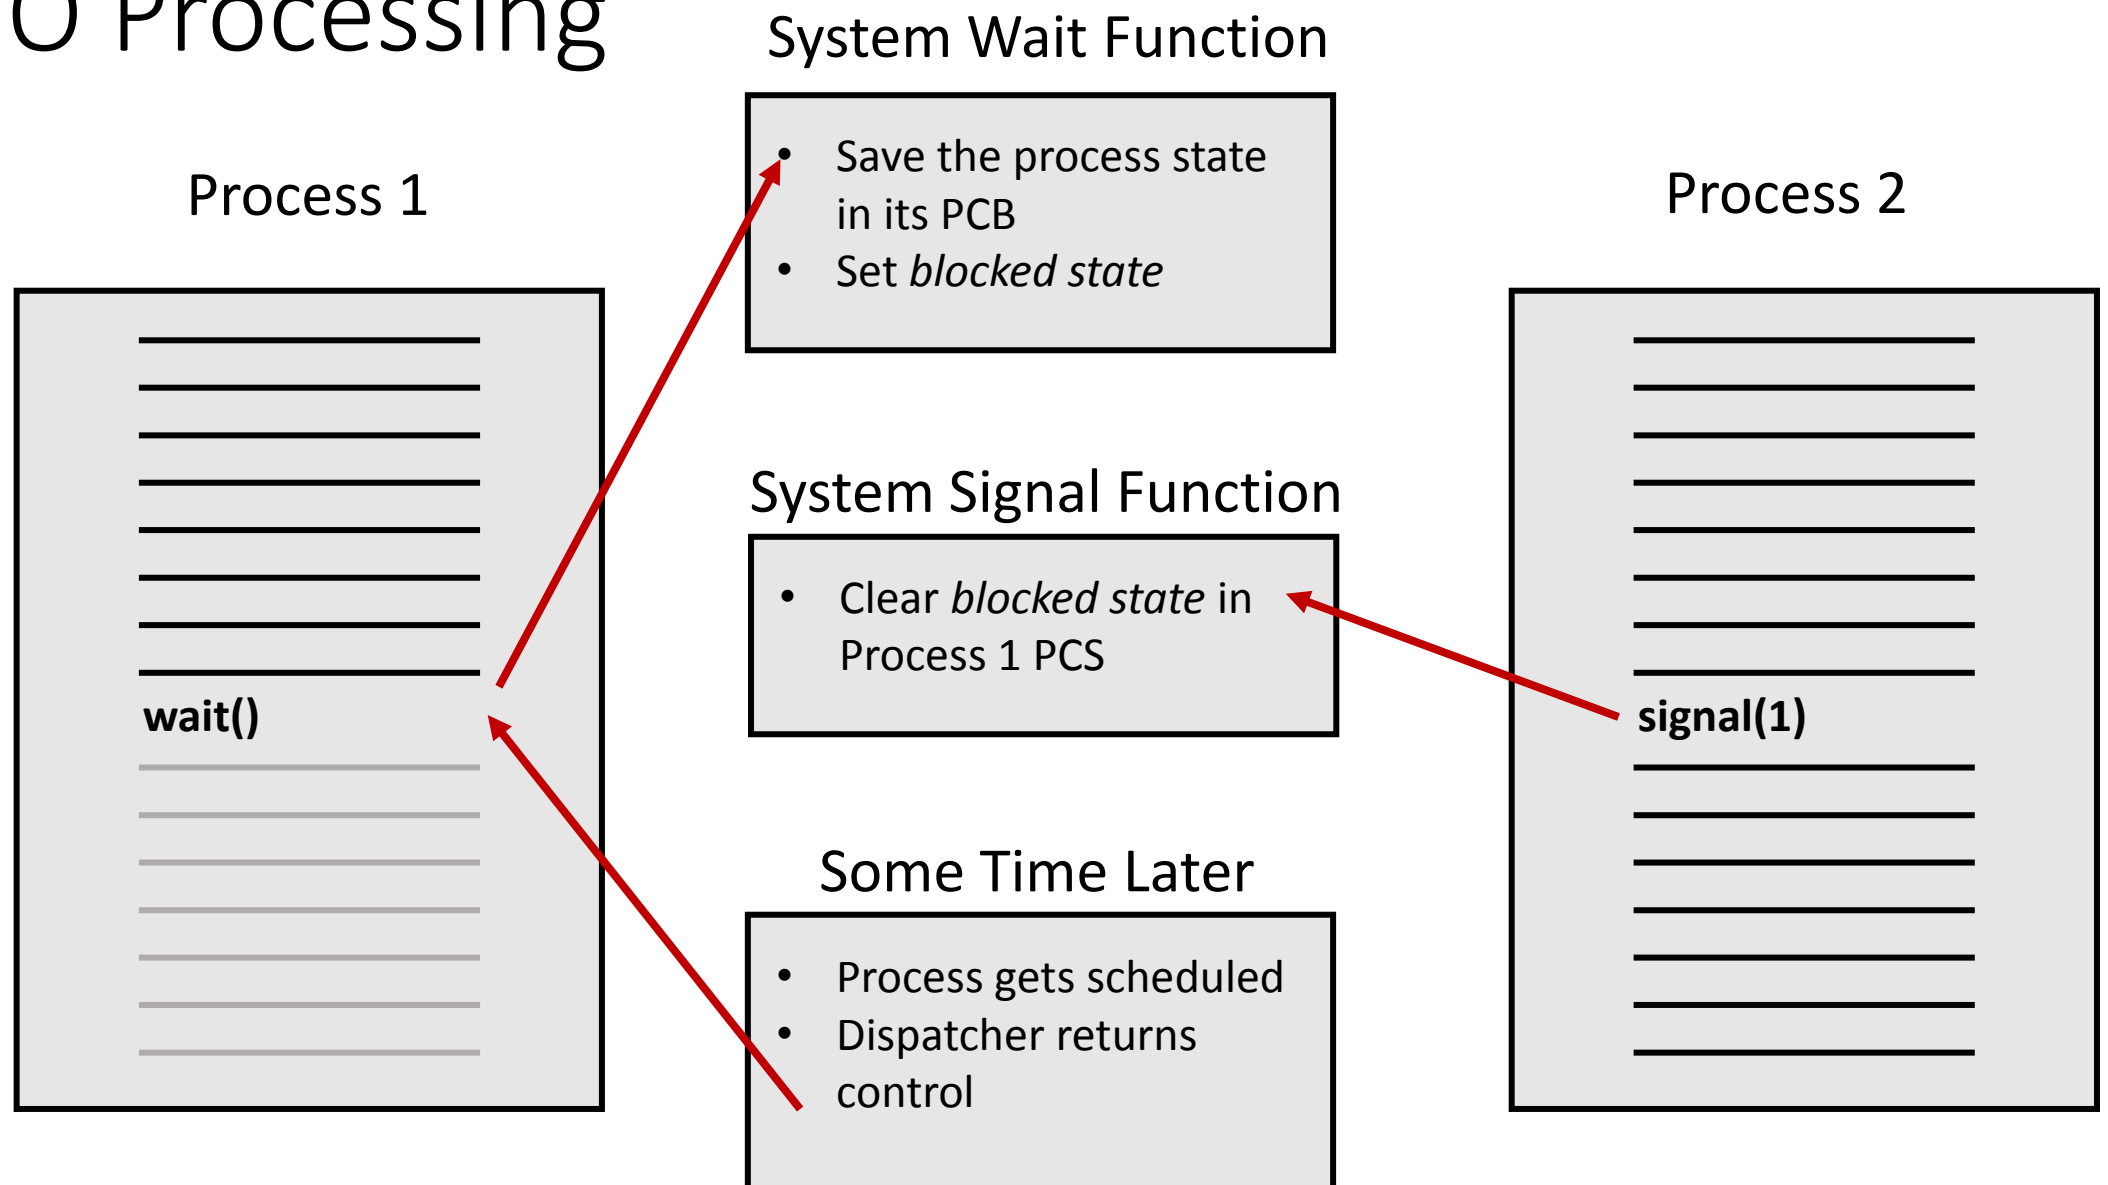
I/O Processing



Process Synchronization and Communication

- A process can pause, waiting for the completion of another process using the *wait()* function call
 - A POSIX shell may launch a child process to run a program from the command line, and wait for its completion
- **Inter-Process Communication (IPC):** A process may make a system call to send a signal to another process, or to set up a signal handler to receive signals
- Just as with I/O wait states, these waits are indicated by a special *blocked* status in the PCB, preventing the process from being scheduled until the condition is met

I/O Processing



(2) Input / Output (I/O) Management

- On early computers, I/O was performed by the processor:

WAITING	TD	DEVICE	; wait until device is ready
	JEQ	WAITING	
	RD	DEVICE	; get a byte from device
	STA	BUFFER, X	; store byte in buffer

- Of course, today we don't tie up the processor waiting for each byte!

Adding a Primitive I/O Subsystem

- System functions to READ and WRITE data
 - READ(device, buffer, count);
 - WRITE(device, buffer, count);
 - The calling process is placed in an *I/O wait state*
 - Functions place the parameters in an I/O Control Block (IOCB)
- An *I/O Process*
 - The I/O process contains a loop: for each IOCB:
 - Check to see if the device is ready
 - If so, READ or WRITE the next byte
 - If the I/O operation is complete, set the status to *ready* in the calling process's PCB

A Modern I/O Subsystem

- On modern systems, I/O is handled by hardware, rather than relying on the processor for byte-by-byte transfers
- System functions to READ and WRITE data
 - READ(device, buffer, count);
 - WRITE(device, buffer, count);
 - The calling process is placed in an *I/O wait state*
 - Functions place the parameters in an I/O Control Block (IOCB)
 - Functions initiate the I/O transfer using dedicated I/O controllers

A Modern I/O Subsystem - Continued

- We do not need an *I/O process*
- The I/O controllers will generate an interrupt when the transfer is complete
- The interrupt handler will set the status to *ready* in the calling process's PCB

(3) Memory Management

Physical Addressing

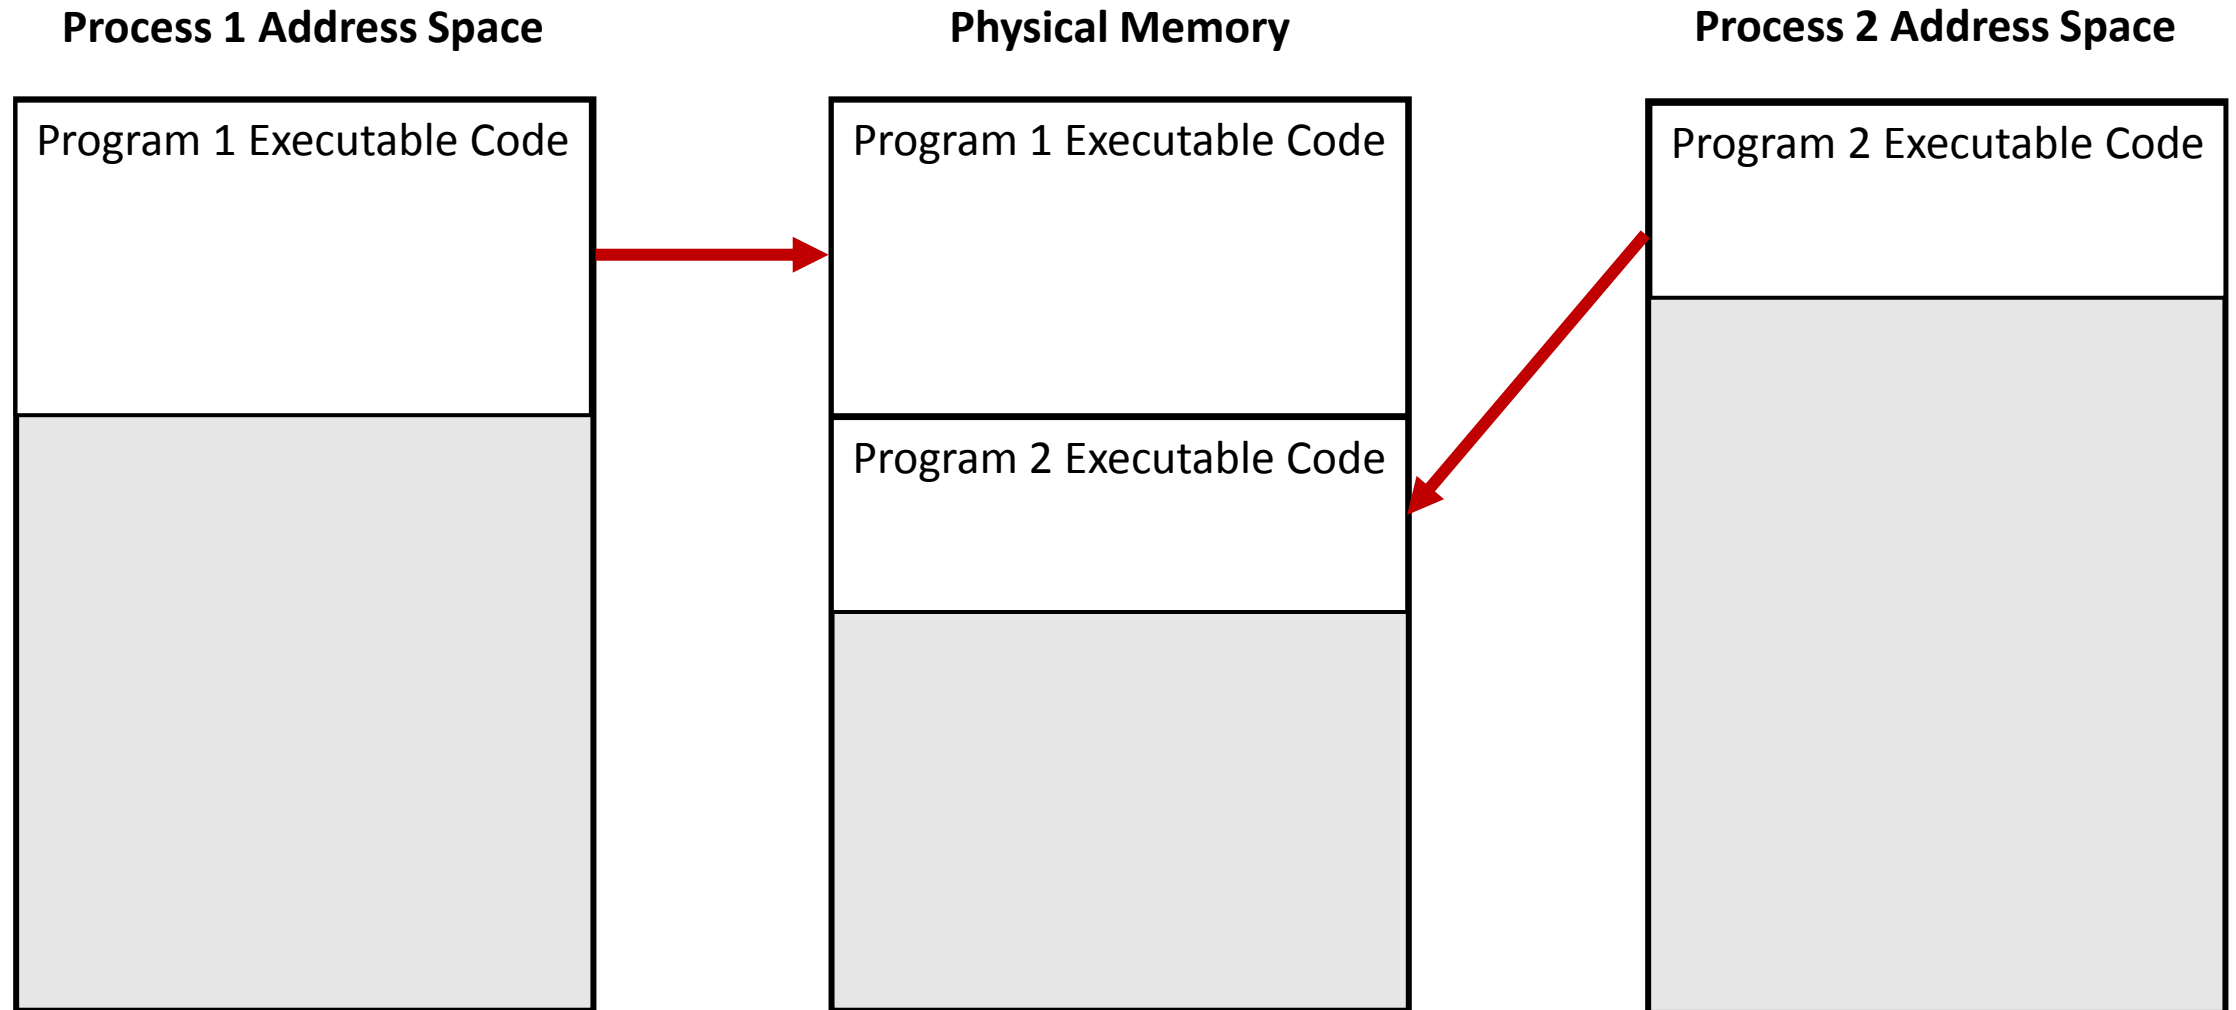
- Early computers did not have any form of address translation
- Addresses used in the program exactly correspond to physical memory addresses
- The size of physical memory was limited to the address range of the instruction set
- To load multiple programs or code blocks into memory, the system required a *relocating loader*
- Microcomputers in the 1970s and embedded computers today use physical addressing

Partitioned Memory

Partitioned Addressing


- Partitioned Addressing requires hardware support – a Memory Management Unit (MMU)
- A “base address” register that can be set when a process is dispatched, and possibly start and end addresses to provide memory protection
- Memory references are translated in hardware by adding the “base address register” to determine the address in physical memory
- The amount of memory available to any process may be limited by the instruction set, but physical memory may be much larger

Partitions




Memory Fragmentation

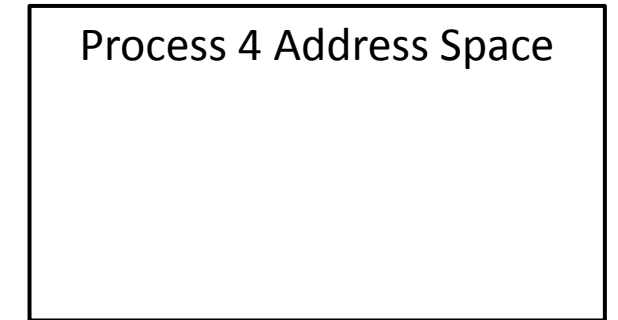
Starting State



Process 2 Ends




Process 4 Doesn't Fit!



We have enough
memory... but it's
not contiguous

Memory Fragmentation – Dynamic Relocation

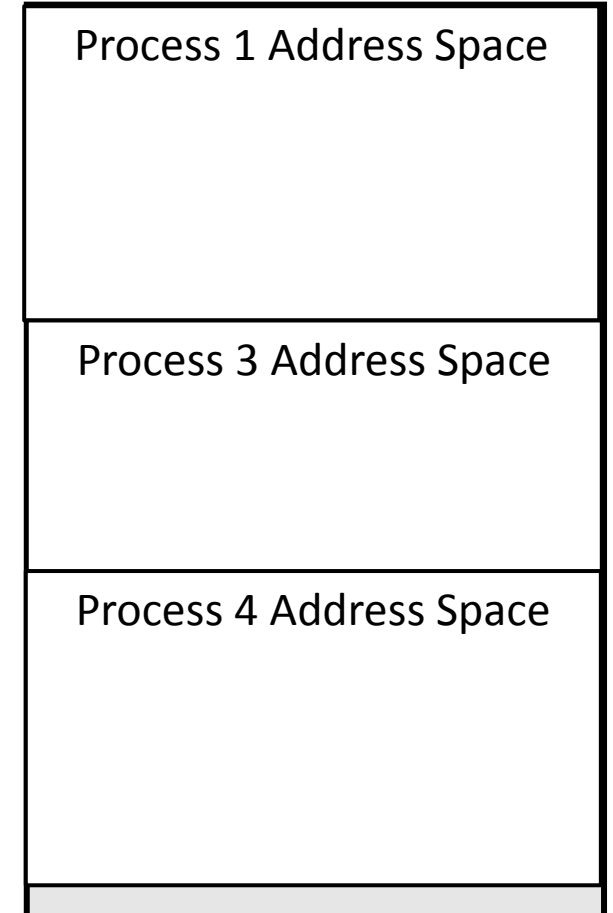
Starting State



Process 2 Ends



Process 4 Starts

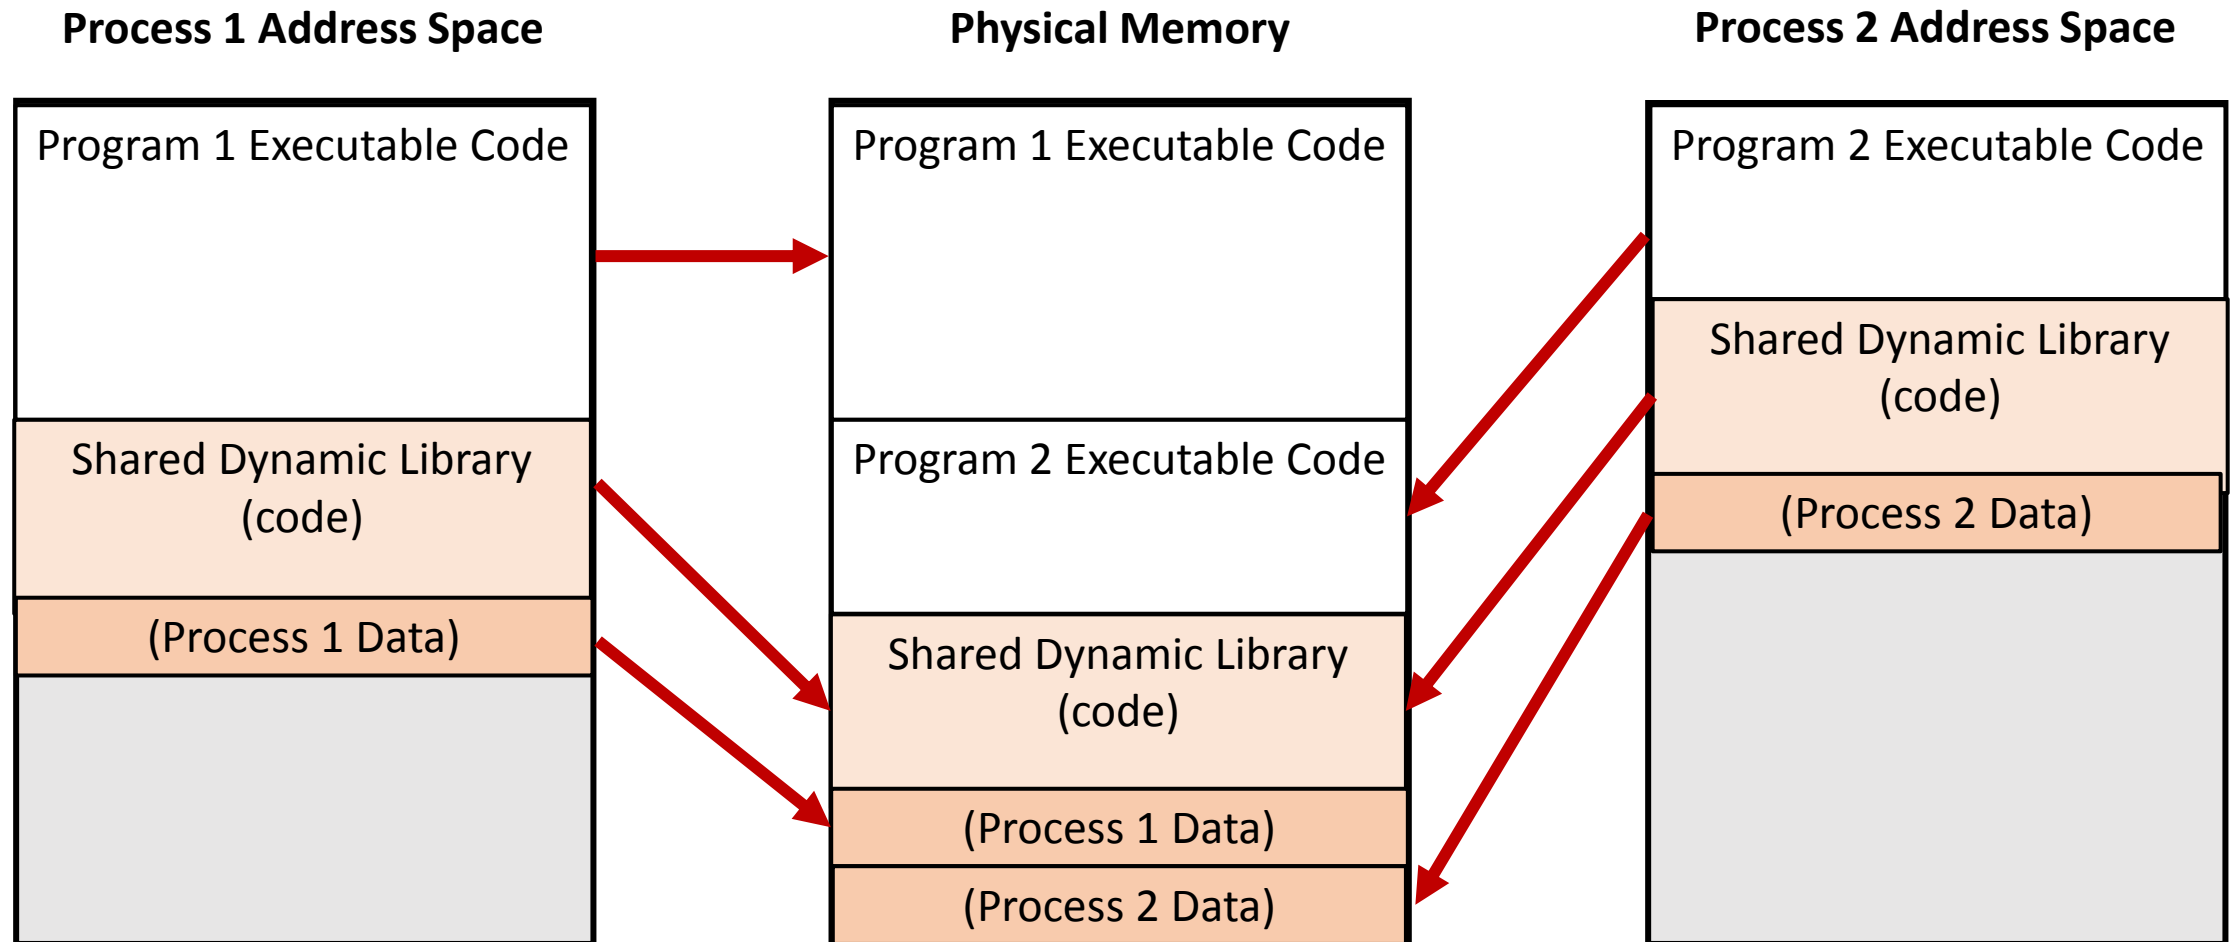


Memory Protection

- *Memory Protection* prevents one process from reading or writing memory used by another process
- In addition to a *base address*, the MMU may support a *start address* and *end address*.
- For each memory reference, the MMU adds the base address, and determines if the resulting physical address live within the bounds of the start and end
 - If not, an error interrupt is generated

Getting Fancy – Modern MMUs

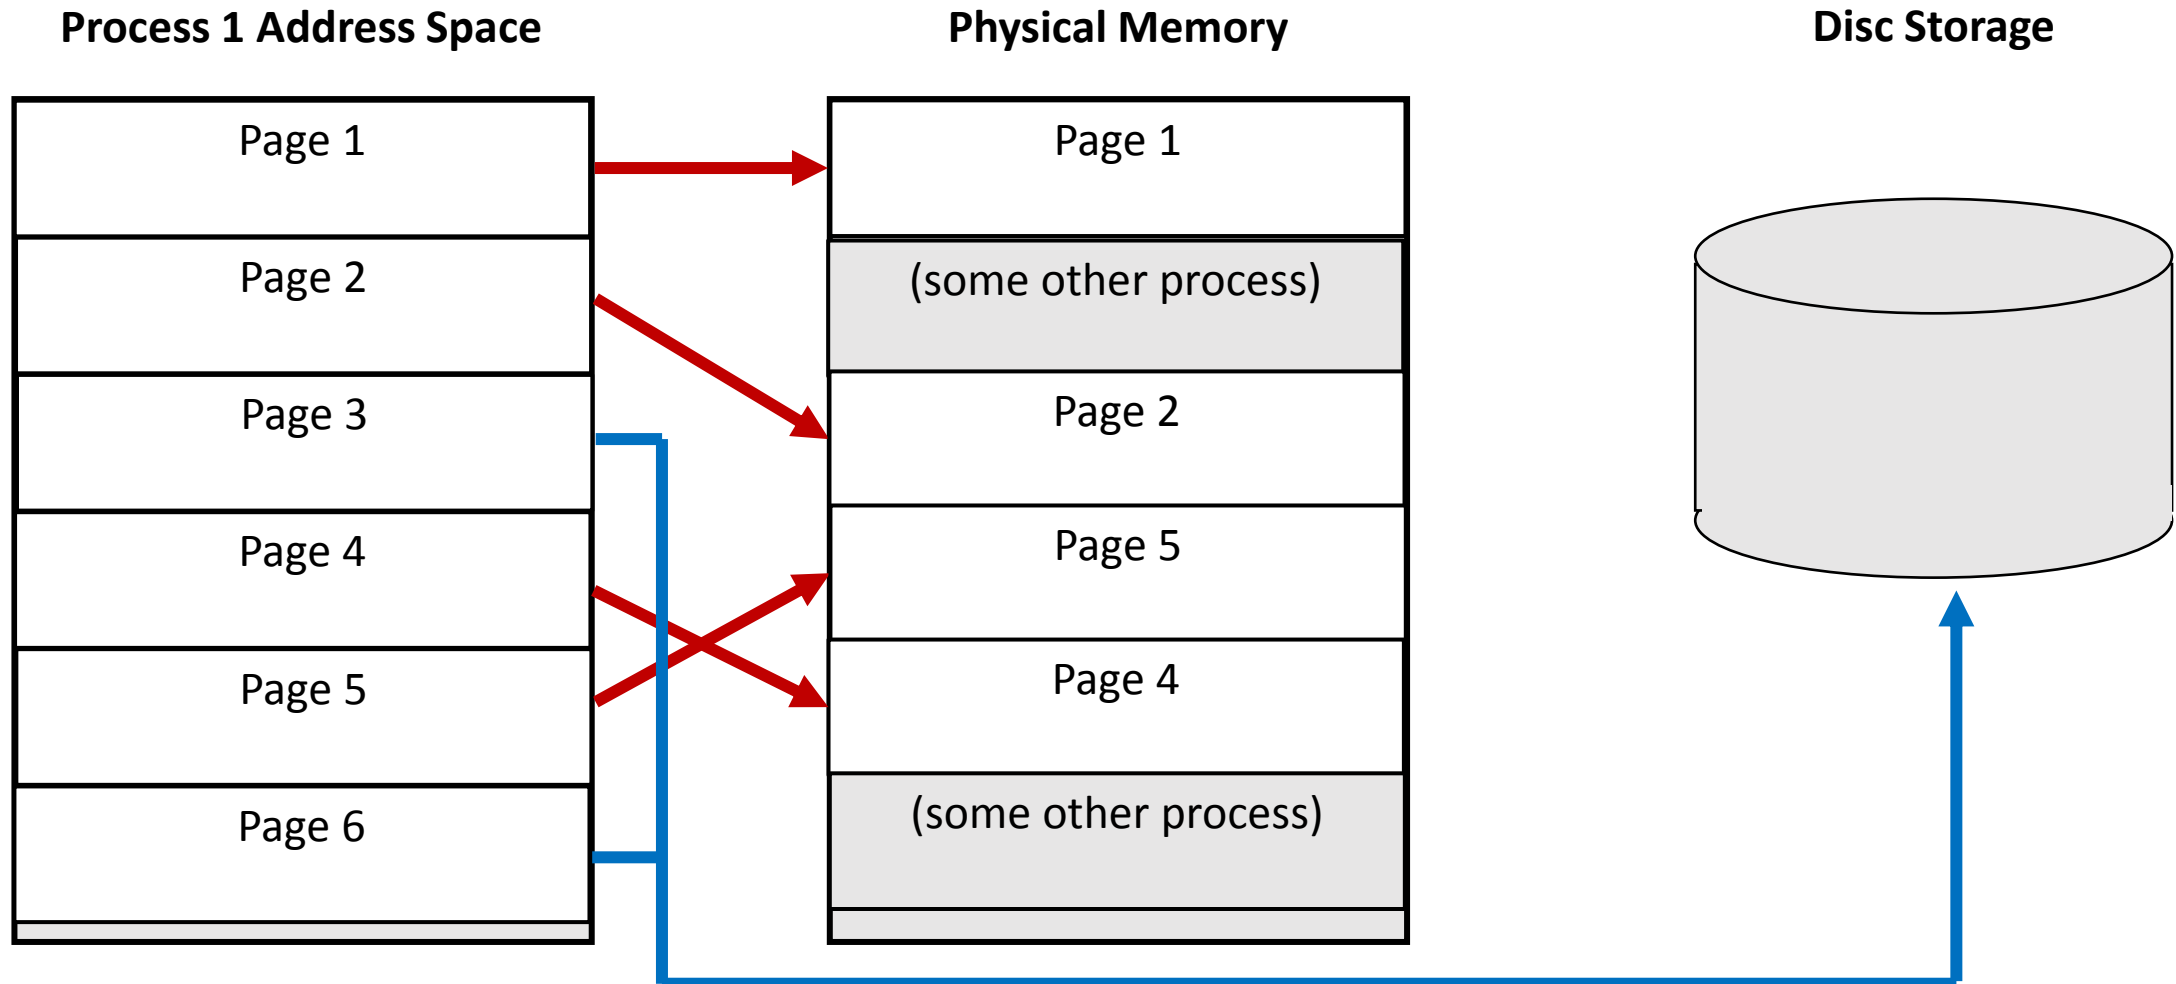
- Separate base addresses for code, data, and shared memory



Virtual Memory

- **Virtual Memory:** a system of software and hardware that allows portions of a program's memory to be temporarily cached on disk, to minimize the requirements for physical memory.
- Allows the operating system to load and execute programs with memory requirements that total far more than the available physical memory.
- Virtually memory systems divide the memory address space into *pages*, each of which is separately mapped by the MMU.
- Not all pages reside in physical memory. Pages that are not loaded in physical memory are cached on disk.

Virtual Memory Mapping



Virtual Memory Mapping

- Requires a separate *base address* for each *page* in the process's address space.
- The MMU maps each memory reference to the corresponding location in physical memory by adding the base address.
- If the page does not reside in physical memory, the MMU will generate a *page fault* interrupt.
 - The process will be placed in an I/O wait state
 - The requested memory page will be loaded from disk into physical memory
 - This may require a page that is currently in physical memory to be moved to disk (swapped out) in order to free up space

Choosing the Page to Swap Out

- If a page is already on disk, and the copy in memory has not been modified, it doesn't actually need to be written to disk – so it's a good candidate to free up
- Other algorithms include:
 - FIFO – First In, First Out
 - LRU – Least Recently Used
 - LFU – Least Frequently Used

Why the Algorithm is Important

- Going to disk to “swap in” a page (and possibly “swap out” a page to free up memory) dramatically slows program execution
- If a RAM access took one SECOND, a disk access would take one DAY
- A system that has insufficient physical memory needs to swap frequently, resulting in *thrashing*, a condition in which the system is essentially unable to perform useful work because of constant page swaps.